

# **Samvaad (Sarvam)<> Vistaar (Ekstep) Integration**

Samvaad, the conversational platform of Sarvam can handle voice-to-voice calls, rich conversations over WhatsApp, and in-app usecases like embedding within a website or mobile applications.

## **Agent Pipeline of Samvaad**

In the context of LLM Orchestration, Samvaad provides the following capabilities:

1. LLM Configuration - Sarvam Medium, Sarvam Tota, Sarvam Pro etc.
2. State management - memory (across conversations), thread context (within conversation).
3. Orchestration - tool calling, ReAct loops, and other agentic behaviour.

## **BYOLO**

To support BYOLO, Samvaad can expose an integration option where users can specify their endpoint configuration for “system of LLMs”, which take care of the above capabilities and Samvaad acts as a channel and normalisation layer to receive user inputs as voice / widgets / text and send the user message to this endpoint. Concretely, the following can be handled by Samvaad:

1. Scheduling (if outbound)
2. Inbound handling with channel integrations
3. ASR + VAD
4. BYOLO instead of Samvaad's agent pipeline
5. Translation/Transliteration/Post LLM as configured by the author
6. TTS
7. Analytics

## Integration

To standardise based on what's already available in the industry, we can go with the following, which is similar to a single LLM with an external inference setup:

1. OpenAI-compatible API
2. Server-side events (SSE)
3. Authentication (Samvaad <> BYOLO)
4. User ID and Tenant ID
5. Payload structure within the `text` response

When you want to say the response to the user:

```
{  
  "audio": "Thank you for confirming."  
}
```

When you need to end the conversation with a sentence:

```
{  
  "end_interaction": true,  
  "audio": "..."  
}
```

Sample code based on OpenAI Chat Completion

```
from fastapi import FastAPI  
from fastapi.responses import StreamingResponse  
from pydantic import BaseModel  
from typing import List, Optional, Literal  
import json  
import time  
  
app = FastAPI()
```

```

class Message(BaseModel):
    role: Literal["system", "user", "assistant"]
    content: str

class ChatCompletionRequest(BaseModel):
    model: str = "gpt-3.5-turbo"
    messages: List[Message]
    stream: Optional[bool] = True
    temperature: Optional[float] = 1.0
    max_tokens: Optional[int] = None

def generate_stream():
    """Generate OpenAI-compatible SSE streaming response"""

    # Static response text that will be streamed
    response_text = "Hello! This is a static streaming response from the Fast
API server."

    # Create a unique ID for this completion
    completion_id = f"chatcmpl-{int(time.time())}"
    created_timestamp = int(time.time())

    # Stream each word as a chunk
    words = response_text.split()

    for i, word in enumerate(words):
        chunk = {
            "id": completion_id,
            "object": "chat.completion.chunk",
            "created": created_timestamp,
            "model": "gpt-3.5-turbo",
            "choices": [
                {
                    "index": 0,
                    "delta": {
                        "content": word + " " if i < len(words) - 1 else word

```

```

        },
        "finish_reason": None
    }
]
}

yield f"data: {json.dumps(chunk)}\n\n"
time.sleep(0.1) # Simulate streaming delay

# Send final chunk with finish_reason
final_chunk = {
    "id": completion_id,
    "object": "chat.completion.chunk",
    "created": created_timestamp,
    "model": "gpt-3.5-turbo",
    "choices": [
        {
            "index": 0,
            "delta": {},
            "finish_reason": "stop"
        }
    ]
}
yield f"data: {json.dumps(final_chunk)}\n\n"

# Send [DONE] message
yield "data: [DONE]\n\n"

```

```

@app.post("/v1/chat/completions")
async def chat_completions(request: ChatCompletionRequest):
    """
    OpenAI-compatible chat completions endpoint with streaming support.
    """

```

This endpoint accepts messages in the standard OpenAI API format and returns a static streamed response.

```
"""

```

```
if request.stream:
    return StreamingResponse(
```

```

        generate_stream(),
        media_type="text/event-stream",
        headers={
            "Cache-Control": "no-cache",
            "Connection": "keep-alive",
        }
    )
else:
    # Non-streaming response (optional)
    return {
        "id": f"chatcmpl-{int(time.time())}",
        "object": "chat.completion",
        "created": int(time.time()),
        "model": request.model,
        "choices": [
            {
                "index": 0,
                "message": {
                    "role": "assistant",
                    "content": "Hello! This is a static response from the FastAPI s
erver."
                },
                "finish_reason": "stop"
            }
        ],
        "usage": {
            "prompt_tokens": 10,
            "completion_tokens": 10,
            "total_tokens": 20
        }
    }

@app.get("/")
async def root():
    return {"message": "OpenAI-compatible API server", "endpoint": "/v1/ch
at/completions"}

```

```
if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000)
```

Reference: <https://platform.openai.com/docs/api-reference/chat/create>

## Out of Scope

1. Language Switch
2. Analytics on Samvaad
- 3.